

# Data Driven Sampling (Lecture 7)

## Classification

Hongwei YUAN



University of Macau

# Classification example: customer's credibility

- Bank observes a customer's yearly income and saving, and determines whether he/she is low-risk or high-risk customer according to the yearly income and saving.
- We can model this problem as:
  - ▶ customer's yearly income and saving is denoted by random variables  $X_1$  and  $X_2$  respectively;
  - ▶ the credibility of customer is modeled by a random variable  $C$  with value 0 or 1, where 1 means high-risk customer and 0 means low-risk customer;
  - ▶ given  $(X_1, X_2)$ , we need to find a formula for  $P(C = 0|X_1, X_2)$  and  $P(C = 1|X_1, X_2)$  to classify customer's credibility, e.g. for any  $(x_1, x_2)$

$$\text{Choose } \begin{cases} 1, & P(C = 1|X_1 = x_1, X_2 = x_2) \geq 0.5, \\ 0, & P(C = 1|X_1 = x_1, X_2 = x_2) < 0.5. \end{cases}$$

# Classification example:



## Two classes: $C_1$ and $C_2$

- The observed data  $\mathbf{x}$  are in  $\mathbb{R}^d$ , and belong to the class  $C_1$  or the class  $C_2$ .
- For the class  $C_1$ , the data density is  $p(\mathbf{x}|C_1)$  for the class  $C_2$ , the data density is  $p(\mathbf{x}|C_2)$
- An example of the above setting:  $C_1$  and  $C_2$  are two car brands,  $x$  is the incomes of car customers.

## Two classes: $C_1$ and $C_2$

- For the class  $C_i$  ( $i=1,2$ ), Likelihood function is

$$L_i(\mathbf{x}) = p(\mathbf{x}|C_i)P(C_i).$$

log likelihood functions is

$$l_i(\mathbf{x}) = \log p(\mathbf{x}|C_i) + \log P(C_i).$$

- Given a data  $\mathbf{x} \in \mathbb{R}^d$ , it is natural to compare  $l_i(\mathbf{x})$  and
  - ▶ Choose  $C_1$  if  $l_1(\mathbf{x}) \geq l_2(\mathbf{x})$ ,
  - ▶ Choose  $C_2$  if  $l_1(\mathbf{x}) < l_2(\mathbf{x})$ ,
  - ▶ The border for classification:  $\{x : l_1(x) = l_2(x)\}$ , it is often a curve or a line.

# Two classes classification: two car brand example

- $C_1$  and  $C_2$  are two car brands,  $x$  is the incomes of car customers.
- For the class  $C_i$  ( $i=1,2$ ), the data density is

$$p(x|C_i) = \frac{1}{\sqrt{2\pi}\sigma_i} e^{-\frac{(x-\mu_i)^2}{2\sigma_i^2}},$$

where  $\mu_i$  and  $\sigma_i$  ( $i=1,2$ ) are all unknown.

- We know

$$l_i(x) = \log P(C_i) - \frac{(x - \mu_i)^2}{2\sigma_i^2} - \frac{1}{2} \log(2\pi) - \log \sigma_i.$$

In order to compare  $l_1(x)$  and  $l_2(x)$ , one needs to know  $P(C_i)$ ,  $\mu_i$  and  $\sigma_i$  ( $i=1,2$ ).

# Two classes classification: two car brand example

- We will estimate  $P(C_i)$ ,  $\mu_i$  and  $\sigma_i$  ( $i=1,2$ ) by observed data. More precisely, suppose that we observed data

$$\{(x_1, r_1), \dots, (x_n, r_n)\}$$

where  $x_i$  is the  $i$ th customer's income and  $r_i$  is the choice of the  $i$ th customer, i.e.  $r_i = C_1$  or  $r_i = C_2$ .

- Denote by  $n_1$  the number of customers who chose  $C_1$ , and by  $n_2$  the number of customers who chose  $C_2$ , we have  $n_1 + n_2 = n$ . we classify the observed data into two classes, one choosing  $C_1$ , the other choosing  $C_2$ , i.e.,

$$\{(x_1^1, C_1), \dots, (x_{n_1}^1, C_1)\}, \quad \{(x_1^2, C_2), \dots, (x_{n_2}^2, C_2)\}$$

# Two classes classification: two car brand example

- We estimate

- ▶  $P(C_1)$  and  $P(C_2)$  by

$$\hat{P}(C_1) = \frac{n_1}{n}, \quad \hat{P}(C_2) = \frac{n_2}{n};$$

- ▶  $\mu_1$  and  $\sigma_1^2$  by

$$\hat{\mu}_1 = \frac{\sum_{i=1}^{n_1} x_i^1}{n_1}, \quad \hat{\sigma}_1^2 = \frac{\sum_{i=1}^{n_1} (x_i^1 - \hat{\mu}_1)^2}{n_1}.$$

- ▶  $\mu_2$  and  $\sigma_2^2$  by

$$\hat{\mu}_2 = \frac{\sum_{i=1}^{n_2} x_i^2}{n_2}, \quad \hat{\sigma}_2^2 = \frac{\sum_{i=1}^{n_2} (x_i^2 - \hat{\mu}_2)^2}{n_2}.$$



# Two classes classification: two car brand example

- Plugging the estimated  $\hat{P}(C_i)$ ,  $\hat{\mu}_i$  and  $\hat{\sigma}_i$  ( $i=1,2$ ) into the log likelihood function  $l_i(x)$  ( $i=1,2$ ) to obtain an estimated log likelihood function:

$$\hat{l}_i(x) = \log \hat{P}(C_i) - \frac{(x - \hat{\mu}_i)^2}{2\hat{\sigma}_i^2} - \frac{1}{2} \log(2\pi) - \log \hat{\sigma}_i.$$

- Now give a new data  $x \in \mathbb{R}$ , we can compute its  $\hat{l}_1(x)$  and  $\hat{l}_2(x)$ , and determine its class by
  - ▶ if  $\hat{l}_1(x) \geq \hat{l}_2(x)$ , we claim  $x$  belongs to  $C_1$ ,
  - ▶ if  $\hat{l}_1(x) < \hat{l}_2(x)$ , we claim  $x$  belongs to  $C_2$ ,
  - ▶ the border is  $\{x : \hat{l}_1(x) = \hat{l}_2(x)\}$ , which is often a curve or a line.

# General framework for the classification problem

- What we can learn from the previous example:
  - ▶ We have seen from the previous slides, we take the two likelihood function  $l_1(x)$  and  $l_2(x)$  as a criterion to determine which class a given data belongs to.
  - ▶  $l_1(x)$  and  $l_2(x)$  are called discriminant functions in classification problem.
- General framework for classification problem:
  - ▶ K classes:  $C_1, \dots, C_K$ ,
  - ▶ K discriminant functions:  $g_1(x), \dots, g_K(x)$ ,
  - ▶ Criterion: for a given data  $x$ , if  $g_k(x) = \max_{1 \leq i \leq K} g_i(x)$ , we choose  $C_k$  as the class of  $x$ .
  - ▶ We will have  $\frac{K(K-1)}{2}$  borders, i.e.  $\{x : g_i(x) = g_j(x)\}$  for  $i \neq j$ .

# Linear discriminant

We still consider two classes classification problem. Let the two discriminant functions be  $g_1(\mathbf{x})$  and  $g_2(\mathbf{x})$  with  $\mathbf{x} \in \mathbb{R}^d$ . Suppose that

$$g_1(\mathbf{x}) = \theta_1^T \mathbf{x} + \theta_{1,0}, \quad g_2(\mathbf{x}) = \theta_2^T \mathbf{x} + \theta_{2,0},$$

where  $\theta_i = (\theta_{i,1}, \dots, \theta_{i,d})^T$  for  $i = 1, 2$ .

- If  $g_1(\mathbf{x}) \geq g_2(\mathbf{x})$  for  $\mathbf{x}$ , we choose the class  $C_1$  for  $\mathbf{x}$ . Otherwise, we choose  $C_2$  for  $\mathbf{x}$ .
- Define  $g(\mathbf{x}) = g_1(\mathbf{x}) - g_2(\mathbf{x})$ , then we have
  - ▶ If  $g(\mathbf{x}) > 0$ , choose  $C_1$ ; if  $g(\mathbf{x}) \leq 0$ , choose  $C_2$ .
  - ▶ Since  $g(\mathbf{x})$  is linear, we use a plane  $g(\mathbf{x}) = 0$  to separate the two classes.

# A simple example for linear discriminant

We have a random number  $\mathbf{z}$  which is either from  $N(\mu_1, \Sigma_1)$  or from  $N(\mu_2, \Sigma_2)$ , with (here we assume that  $P(C_1) = P(C_2) = 1/2$ )

$$\mu_1 = [2, 2]^T, \quad \mu_2 = [0, 0]^T, \quad \Sigma_1 = \Sigma_2 = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

- **Assignment 1:** Compute log likelihood function  $l_1(\mathbf{x}) = -\frac{(x_1-2)^2 + (x_2-2)^2}{2} + C$  and  $l_2(\mathbf{x}) = -\frac{x_1^2 + x_2^2}{2} + C$  with  $C = -\log(4\pi)$ .
- As shown in the previous slide, we take  $l_1(\mathbf{x})$  and  $l_2(\mathbf{x})$  as the two discriminant functions. Denote  $l(\mathbf{x}) = l_1(\mathbf{x}) - l_2(\mathbf{x}) = 2x_1 + 2x_2 - 4$ .
- If  $l(\mathbf{z}) > 0$ ,  $\mathbf{z}$  belongs to class 1. If  $l(\mathbf{z}) \leq 0$ ,  $\mathbf{z}$  belongs to class 2.

# A simple example for linear discriminant

Question: I don't know the sources of these data, but only know their covariances are both identity matrix. How can I classify a given data  $\mathbf{z}$ ?

- You sample  $n_1$  random numbers  $\mathbf{x}_1, \dots, \mathbf{x}_{n_1}$  from  $N(\mu_1, \Sigma_1)$  and another  $n_2 = 30$  random numbers  $\mathbf{y}_1, \dots, \mathbf{y}_{n_2}$  from  $N(\mu_2, \Sigma_2)$  for me.
- I use these data to estimate  $\mu_1$  and  $\mu_2$  by

$$\hat{\mu}_1 = \frac{\mathbf{x}_1 + \dots + \mathbf{x}_{n_1}}{n_1}, \hat{\mu}_2 = \frac{\mathbf{y}_1 + \dots + \mathbf{y}_{n_2}}{n_2}$$


- Replacing  $\mu_1, \mu_2$  by  $\hat{\mu}_1, \hat{\mu}_2$  respectively, we get  $\hat{l}_1(\mathbf{x}), \hat{l}_2(\mathbf{x})$  and  $\hat{l}(\mathbf{x})$ .
- (Class Exercise) Show that  $\hat{l}(\mathbf{x}) = (\hat{\mu}_1 - \hat{\mu}_2)^T \mathbf{x} + \frac{1}{2}((\hat{\mu}_2)^T \hat{\mu}_2 - (\hat{\mu}_1)^T \hat{\mu}_1)$ .

```

import numpy as np
from scipy.stats import cauchy
from scipy.optimize import minimize_scalar
import matplotlib.pyplot as plt
from sympy import *

num1,num2=300,300
np.random.seed(0)
mean_1 = [0, 0]
cov_1 = [[0.1, 0], [0, 0.1]]
mean_2 = [2, 2]
cov_2 = [[0.1, 0], [0, 0.1]]
x1=np.random.multivariate_normal(mean_1, cov_1, num1)
x2=np.random.multivariate_normal(mean_2, cov_2, num2)
plt.scatter(x2[:,0],x2[:,1],marker='o')
plt.scatter(x1[:,0],x1[:,1],marker='x')
#def p(a,b):
#    return a/(a+b)
average_mu1=sum(x1)/num1
average_mu2=sum(x2)/num2
print(average_mu1)
print(average_mu2)
def l1(x):
    return np.dot((x-average_mu1),(x-average_mu1))
def l2(x):
    return np.dot((x-average_mu2),(x-average_mu2))
z_test=(4,5)
if l1(z_test)>l2(z_test):
    print('z_test belongs to class 1')
else: print('z_test belongs to class 2')
x=np.arange(-3,3,0.1)
a=(average_mu2[0]-average_mu1[0])/(average_mu2[1]-average_mu1[1])
b=4/(average_mu2[1]-average_mu1[1])
print(a)
print(b)
y=b-a*x
#print(y)
plt.plot(x,y)
plt.savefig('class.pdf')
plt.show()

```



```


import numpy as np
import matplotlib.pyplot as plt
from sklearn import svm

mean1 = (1,1)
mean2 = (2,2)
cov1 = np.array([[0.5,0],[0,0.5]])
cov2 = np.array([[1,0],[0,1]])
np.random.seed(10)
x1 = np.random.multivariate_normal(mean1, cov1, (50,))
x2 = np.random.multivariate_normal(mean2, cov2, (50,))
x = np.r_[x1, x2]
y = [0]*50+[1]*50 #class labels

#SVM
clf = svm.SVC(kernel = 'linear')
clf.fit(x, y)
#build the boundary
w = clf.coef_[0]
xx = np.linspace(0, 7)
a = -w[0] / w[1]
yy = a*xx - (clf.intercept_ / w[1])
plt.scatter(x1[:, 0], x1[:, 1])
plt.scatter(x2[:, 0], x2[:, 1])
plt.plot(xx, yy, 'k-')
plt.savefig("classification graph.pdf")
plt.show()

```





# Logistic regression (1 dimension case)

Let  $X$  be a continuous random variable (valued on  $\mathbb{R}$ ), which can be in the class  $C_1$  or  $C_2$  with probability  $P(C_1)$  and  $P(C_2)$  respectively.

- Let  $p(x|C_i)$  be conditional probability density of  $X$  given  $C_i$  ( $i = 1, 2$ ).
- Let  $p(x)$  be the probability density of  $X$ , we have

$$p(x) = p(x|C_1)P(C_1) + p(x|C_2)P(C_2).$$

- Let  $P(C_i|X = x)$  be conditional probability of  $C_i$  given  $X = x$ , we simply denote it by  $P(C_i|x)$  ( $i = 1, 2$ ).
- Let  $p(C_i, x)$  be the probability density of  $x$  and  $C_i$  ( $i = 1, 2$ ).

# Logistic regression (1 dimension case)

- One assumes

$$\log \frac{p(x|C_2)}{p(x|C_1)} = \beta_1 x + \beta_0,$$

where  $\beta_0$  and  $\beta_1$  is some unknown parameter. By Bayes' rule,

$$\frac{P(C_2|x)}{P(C_1|x)} = \frac{p(C_2, x)/p(x)}{p(C_1, x)/p(x)} = \frac{p(x|C_2)P(C_2)}{p(x|C_1)P(C_1)},$$

which implies

$$\log \frac{P(C_2|x)}{P(C_1|x)} = \beta_1 x + \beta'_0, \quad (*)$$

where  $\beta'_0 = \beta_0 + \log \frac{P(C_2)}{P(C_1)}$ .

- By the above deviation, we assume

$$\log \frac{P(C_2|x)}{P(C_1|x)} = \beta'_0 + \beta_1 x,$$

it is called **logistic regression model** with parameters  $\beta'_0, \beta_1$ .

# Logistic regression (1 dimension case)

- We have

$$P(C_2|x) + P(C_1|x) = 1. \quad (**)$$

- By (\*) and (\*\*),

$$P(C_2|x) = \frac{e^{\beta'_0 + \beta_1 x}}{1 + e^{\beta'_0 + \beta_1 x}} = \frac{1}{1 + e^{-\beta'_0 - \beta_1 x}}, \quad P(C_1|x) = \frac{1}{1 + e^{\beta'_0 + \beta_1 x}}$$

- In Statistics, for two classes classification problem, we often use 1 and 0 to denote classes. For instance,

$$Y = \begin{cases} 1, & \text{if } C_2; \\ 0, & \text{if } C_1. \end{cases}$$

So  $P(Y = 1) = P(C_2)$  and  $P(Y = 0) = P(C_1)$ .

# Logistic regression (1 dimension case)

- We can see  $P(1|x) = \frac{1}{1+e^{-\beta'_0-\beta_1x}}$ ,  $P(0|x) = \frac{1}{1+e^{\beta'_0+\beta_1x}}$ .
- **Assignment 2:** Show that

$$P(y|x) = \frac{1}{1 + e^{(1-2y)(\beta'_0+\beta_1x)}}, \quad y = 0, 1.$$

# Logistic regression (1 dimension case)

Let  $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$  be a sequence of observed data.

- Let the probability of  $X$  is  $p(x)$  (we shall see that the regression does not depends on the probability of  $X$ ), then we have

$$p(1, x) = P(1|x)p(x), \quad p(0, x) = P(0|x)p(x).$$

- The likelihood function of these data are

$$L_n(\beta'_0, \beta_1) = \prod_{i=1}^n p(y_i, x_i) = \prod_{i=1}^n P(y_i|x_i) \prod_{i=1}^n p(x_i) \quad (1)$$

- Assignment 2:** Show that

$$L_n(\beta'_0, \beta_1) = \prod_{i=1}^n \frac{1}{1 + e^{(1-2y_i)(\beta'_0 + \beta_1 x_i)}} \prod_{i=1}^n p(x_i)$$

# Logistic regression (1 dimension case)

- The log likelihood function is

$$l_n(\beta'_0, \beta_1) = \sum_{i=1}^n \log \frac{1}{1 + e^{(1-2y_i)(\beta'_0 + \beta_1 x_i)}} + \sum_{i=1}^n \log p(x_i).$$

- The MLE is

$$\begin{aligned}(\hat{\beta}'_{0,n}, \hat{\beta}_{1,n}) &= \operatorname{argmax}_{(\beta'_0, \beta_1) \in \mathbb{R}^2} l_n(\beta'_0, \beta_1) \\&= \operatorname{argmax}_{(\beta'_0, \beta_1) \in \mathbb{R}^2} \sum_{i=1}^n \log \frac{1}{1 + e^{(1-2y_i)(\beta'_0 + \beta_1 x_i)}} \quad (2) \\&= \operatorname{argmin}_{(\beta'_0, \beta_1) \in \mathbb{R}^2} \sum_{i=1}^n \log(1 + e^{(1-2y_i)(\beta'_0 + \beta_1 x_i)})\end{aligned}$$

# Logistic regression (1 dimension case)

- (Class Exercise)  $(\hat{\beta}'_{0,n}, \hat{\beta}_{1,n})$  satisfies the following equations:

$$\sum_{i=1}^n \frac{1 - 2y_i}{1 + e^{-(1-2y_i)(\beta'_0 + \beta_1 x_i)}} = 0, \quad \sum_{i=1}^n \frac{(1 - 2y_i)x_i}{1 + e^{-(1-2y_i)(\beta'_0 + \beta_1 x_i)}} = 0$$

- Unfortunately, it is hard to solve these equations. We shall learn (stochastic) gradient descent algorithm to solve it later.
- One simplified way is:
  - ▶ Find estimates  $\hat{P}(1|x)$  and  $\hat{P}(0|x)$  of  $P(1|x)$  and  $P(0|x)$  for a given  $x$ .
  - ▶ Consider  $\log \frac{P(1|x_i)}{P(0|x_i)} = \beta'_0 + \beta_1 x_i$  for  $i = 1, \dots, n$ , then we can estimate  $\beta'_0$  and  $\beta_1$  by least square regression.



# Logistic regression (multidimensional case)

Let  $X$  be a random variable (valued on  $\mathbb{R}^d$ ), and let  $Y$  be its label with a value 0 or 1.

- Denote  $P(1|x) = P(Y = 1|X = x)$  and  $P(0|x) = P(Y = 0|X = x)$  for a given  $x = (x_1, \dots, x_d) \in \mathbb{R}^d$ , then

$$P(0|x) = 1 - P(1|x).$$

- (Class Exercise) We shall take

$$P(1|x) = \frac{1}{1 + e^{-\beta'_0 - \beta_1 x_1 - \dots - \beta_d x_d}},$$

$$P(0|x) = \frac{1}{1 + e^{\beta'_0 + \beta_1 x_1 + \dots + \beta_d x_d}}.$$

- We will not enter the details of multidimensional logistic regression.

# A case study of Stock market data (SP 500 stock index)

In this case study, we will use the SP 500 stock index from 2001 to 2005 to predict the direction of stock ('up' or 'down'). The data structure is as the following:

- 1250 lines data, every line has: year, day-5, ..., day-1, volume, today, direction
- (day-5,...,day-1) are the change in SP 500 indexes of the previous 5 days, volume is the shares traded in the previous day (in billion unit), today is today's SP 500 index, direction is today's direction (up or down).
- We shall use the logistic model to fit these data as the following:  
direction of today  $\sim$  day-5,...,day-1, volume.
- We compare the direction predicted by the model with the known direction of today, to evaluate the model.

	Year	day -5	day -4	day -3	day -2	day -1	Volume	Today	Direction	
1	2001	0.381	-0.192	-2.624	-1.055	5.01	1.1913	0.959	Up	
2	2001	0.959	0.381	-0.192	-2.624	-1.055	1.2965	1.032	Up	
3	2001	1.032	0.959	0.381	-0.192	-2.624	1.4112	-0.623	Down	
4	2001	-0.623	1.032	0.959	0.381	-0.192	1.276	0.614	Up	
5	2001	0.614	-0.623	1.032	0.959	0.381	1.2057	0.213	Up	
6	2001	0.213	0.614	-0.623	1.032	0.959	1.3491	1.392	Up	
7	2001	1.392	0.213	0.614	-0.623	1.032	1.445	-0.403	Down	
8	2001	-0.403	1.392	0.213	0.614	-0.623	1.4078	0.027	Up	
9	2001	0.027	-0.403	1.392	0.213	0.614	1.164	1.303	Up	
10	2001	1.303	0.027	-0.403	1.392	0.213	1.2326	0.287	Up	
11	2001	0.287	1.303	0.027	-0.403	1.392	1.309	-0.498	Down	
12	2001	-0.498	0.287	1.303	0.027	-0.403	1.258	-0.189	Down	
13	2001	-0.189	-0.498	0.287	1.303	0.027	1.098	0.68	Up	
14	2001	0.68	-0.189	-0.498	0.287	1.303	1.0531	0.701	Up	
15	2001	0.701	0.68	-0.189	-0.498	0.287	1.1498	-0.562	Down	
16	2001	-0.562	0.701	0.68	-0.189	-0.498	1.2953	0.546	Up	
17	2001	0.546	-0.562	0.701	0.68	-0.189	1.1188	-1.747	Down	
18	2001	-1.747	0.546	-0.562	0.701	0.68	1.0484	0.359	Up	
19	2001	0.359	-1.747	0.546	-0.562	0.701	1.013	-0.151	Down	
20	2001	-0.151	0.359	-1.747	0.546	-0.562	1.0596	-0.841	Down	
21	2001	-0.841	-0.151	0.359	-1.747	0.546	1.1583	-0.623	Down	
22	2001	-0.623	-0.841	-0.151	0.359	-1.747	1.1072	-1.334	Down	
23	2001	-1.334	-0.623	-0.841	-0.151	0.359	1.0755	1.183	Up	
24	2001	1.183	-1.334	-0.623	-0.841	-0.151	1.0391	-0.865	Down	
25	2001	-0.865	1.183	-1.334	-0.623	-0.841	1.0752	-0.218	Down	

# A case study of Stock market data (SP 500 stock index)

We denote  $\mathbf{x} = (\text{day} - 1, \dots, \text{day} - 5, \text{volume})$

- Denote  $y = 0$  if 'up' and  $y = 1$  if 'down', for the  $i$ -th line's data  $\mathbf{x}_i$ ,

$$\hat{P}(0|\mathbf{x}_i) = \frac{\#('+' \text{ index of 5 days})}{5}, \quad \hat{P}(1|\mathbf{x}_i) = \frac{\#('-' \text{ index of 5 days})}{5},$$

- The logistic model is

$$\log \frac{\hat{P}(1|\mathbf{x}_i)}{\hat{P}(0|\mathbf{x}_i)} = \beta^T \mathbf{x}_i + \beta_0$$

where  $\beta = (\beta_6, \beta_5, \dots, \beta_1)$ .

# A case study of Stock market data (SP 500 stock index)

The data of the first line: 5 days' indices=(0.381,-0.192,-2.624,-1.055,5.01), volume=1.1913, i.e.  $\mathbf{x}_1 = (0.381, -0.192, -2.624, -1.055, 5.01, 1.1913)$ .



$$\hat{P}(0|\mathbf{x}_1) = \frac{\#('+' \text{ index of 5 days})}{5} = \frac{2}{5},$$
$$\hat{P}(1|\mathbf{x}_1) = \frac{\#('-' \text{ index of 5 days})}{5} = \frac{3}{5}.$$

- The logistic model for the data of the first line data is

$$\log \frac{3}{2} = \beta^T \mathbf{x}_1 + \beta_0.$$

where  $\beta = (\beta_6, \beta_5, \dots, \beta_1)$  and  $\beta_0$  are coefficients not known.

- Similarly, we can write down 1250 relations as above by all the 1250 lines' data.

# A case study of Stock market data (SP 500 stock index)

- Running the glm(Generalized Linear Model) model function in Python, we can get the values of  $\hat{\beta} = (\hat{\beta}_6, \dots, \hat{\beta}_1, \hat{\beta}_0)$  from the 1250 relations.
- We need to evaluate whether this glm model is good or bad for fitting the SP 500 data, in the following way: for each line's data  $\mathbf{x}_i$

- ▶ Compute

$$a_i = \hat{\beta}^T \mathbf{x}_i + \hat{\beta}_0,$$

if  $a_i > 0$ , we take  $y_i = 1$ , 'down'; if  $a_i \leq 0$ , we take  $y_i = 0$ , 'up'.

- ▶ Compare the above decision ('up' or 'down') with the known 'up' or 'down' in the  $i$ -th line. If they are the same, the prediction is correct. Otherwise, the prediction is wrong.

# A case study of Stock market data (SP 500 stock index)

- Evaluation score =  $\frac{\text{correct predictions}}{1250}$ .
- If the evaluation score is not high, we can conclude that the model is bad for fitting the data.

# A case study of Stock market data (SP 500 stock index)

```
import pandas as pd # pandas includes data structure
import numpy as np
import matplotlib.pyplot as plt
import statsmodels.formula.api as smf # glm is in the package statsmodels
import statsmodels.api as sm
from sklearn.metrics import confusion_matrix # evaluation of the model, diagonal is correct and off-diagonal is wrong
from pandas import Series, DataFrame

data = pd.read_csv("Smarket.csv", index_col=0)
print(data.columns.values.tolist()) #print the colnames
print(data.shape) #print the size of data
print(data.describe()) #descriptive statistical analysis
data1 = data.drop('Direction', axis=1) #delete the qualitative variable 'Direction'
print(data1)
print(data1.corr()) #compute the correlations among the predictors in a data set
plt.plot(data[['Volume']])
plt.show()

# Using the Generalized Linear Model (GLM) to fit the Logistic regression
model = smf.glm("Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 + Volume", family = sm.families.Binomial(), data=data)
result = model.fit()
print(result.summary())
print(np.column_stack((data[["Direction"]].values.flatten(), result.model.endog))) # check the dummy variables !!!down=1

# access the coefficients and p-values
print("Coefficients")
print(result.params)
print()
print("p-Values") # p-values for the coefficients
print(result.pvalues)
print()
# predict the probability the market will go down
print("the probability of market down")
probs = result.predict()
print(probs[1:10])
print()
print("Dependent variables")
print(result.model.endog_names)
print()
```



# A case study of Stock market data (SP 500 stock index)

```
# create a vector of class predictions based on whether the predicted probability of a market
# increase is greater than or less than 0.5
pred = ['Up' if x < 0.5 else 'Down' for x in probs] # definition

# produce a confusion matrix in order to determine how many
# observations were correctly or incorrectly classified
print('confusion matrix')
conmat = confusion_matrix(data["Direction"], pred)
print(conmat)
print()

print('correctly predicted probability')
print(sum(np.diag(conmat))/len(data['Direction']))
print()

print('the fraction of days for which the prediction was correct')
print(np.mean(pred == data['Direction']))#the fraction of days for which the prediction was correct
```